Bootcamp Python

42 ARTIFICIAL INTELLIGENCE

Day04
Pandas

Bootcamp Python

Day04 - Pandas

Today you will learn how to use a Python library that will allow you to manipulate dataframes.

Notions of the day

Pandas! And Bamboos!

General rules

- Use the Pandas Library.

- The version of Python to use is 3.7, you can check the version of Python with the following command: `python -V`

- The norm: during this bootcamp you will follow the PEP 8 standards. You can install pycodestyle which is a tool to check your Python code.

- The function eval is never allowed.

- The exercises are ordered from the easiest to the hardest.

- Your exercises are going to be evaluated by someone else, so make sure that your variable names and function names are appropriate and civil.

- Your manual is the internet.

- You can also ask questions in the dedicated channel in the 42 AI Slack: 42-ai.slack.com.

- If you find any issue or mistakes in the subject please create an issue on our dedicated repository on Github.

Helper

For this day you will use the dataset `athlete_events.csv` provided in the `resources` folder.

```
pip install pandas
```

Ensure that you have the right Python version.

```
> which python
/goinfre/miniconda/bin/python
> python -V
Python 3.7.*
> which pip
/goinfre/miniconda/bin/pip
```

Exercise 00 - FileLoader

Turn-in directory :	ex00
Files to turn in :	FileLoader.py
Allowed libraries :	Pandas
Remarks :	Be as lazy as possible. . .

Write a class named `FileLoader` which implements the following methods:

- `load(path)` : takes as an argument the file path of the dataset to load, displays a message specifying the dimensions of the dataset (e.g. 340 x 500) and returns the dataset loaded as a pandas.DataFrame.

- `display(df, n)` : takes a pandas.DataFrame and an integer as arguments, displays the first n rows of the dataset if n is positive, or the last n rows if n is negative.

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load("../data/adult_data.csv")
Loading dataset of dimensions 32561 x 15
>>> loader.display(data, 12)
age        workclass  fnlwgt  ... hours-per-week  native-country salary
0    39        State-gov   77516  ...             40   United-States  <=50K
1    50  Self-emp-not-inc   83311  ...             13   United-States  <=50K
2    38          Private  215646  ...             40   United-States  <=50K
3    53          Private  234721  ...             40   United-States  <=50K
4    28          Private  338409  ...             40            Cuba  <=50K
5    37          Private  284582  ...             40   United-States  <=50K
6    49          Private  160187  ...             16         Jamaica  <=50K
7    52  Self-emp-not-inc  209642  ...             45   United-States   >50K
8    31          Private   45781  ...             50   United-States   >50K
9    42          Private  159449  ...             40   United-States   >50K
10   37          Private  280464  ...             80   United-States   >50K
11   30        State-gov  141297  ...             40           India   >50K

[12 rows x 15 columns]
```

Note: Your terminal may display more columns if the window is wider.

Exercise 01 - YoungestFellah

Turn-in directory :	ex01
Files to turn in :	FileLoader.py, YoungestFellah.py
Allowed libraries :	Pandas
Remarks :	n/a

This exercise uses the following dataset: `athlete_events.csv`

Write a function `youngestFellah` which takes two arguments:

- a pandas.DataFrame which contains the dataset

- an Olympic year

The function returns a dictionary containing the age of the youngest woman and man who took part in the Olympics on that year. The name of the dictionary's keys is up to you, but it must be self-explanatory.

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
Loading dataset of dimensions 271116 x 15
>>> from YoungestFellah import youngestFellah
>>> youngestFellah(data, 2004)
{'f': 13.0, 'm': 14.0}
```

Exercise 02 - ProportionBySport

Turn-in directory :	ex02
Files to turn in :	FileLoader.py, ProportionBySport.py
Allowed libraries :	Pandas
Remarks :	n/a

This exercise uses the dataset `athlete_events.csv`

Write a function **proportionBySport** which takes four arguments:

- a pandas.DataFrame of the dataset

- an olympic year

- a sport

- a gender

The function returns a float corresponding to the proportion (percentage) of participants who played the given sport among the participants of the given gender.

The function answers questions like the following : "What was the percentage of female basketball players among all the female participants of the 2016 Olympics?"

Hint: here and further, if needed, drop duplicated sportspeople to count only unique ones. Beware to call the dropping function at the right moment and with the right parameters, in order not to omit any individuals.

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
Loading dataset of dimensions 271116 x 15
>>> from ProportionBySport import proportionBySport
>>> proportionBySport(data, 2004, 'Tennis', 'F')
0.01935634328358209
```

We assume that we are always using appropriate arguments as input, and thus do not need to handle input errors.

Exercise 3 - HowManyMedals

Turn-in directory :	ex03
Files to turn in :	FileLoader.py, HowManyMedals.py
Allowed libraries :	Pandas
Remarks :	n/a

This exercise uses the following dataset: `athlete_events.csv`

Write a function `howManyMedals` which takes two arguments:

- a pandas.DataFrame which contains the dataset

- a participant name

The function returns a dictionary of dictionaries giving the number and type of medals for each year during which the participant won medals.
The keys of the main dictionary are the Olympic games years. In each year's dictionary, the keys are 'G', 'S', 'B' corresponding to the type of medals won (gold, silver, bronze). The innermost values correspond to the number of medals of a given type won for a given year.

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
Loading dataset of dimensions 271116 x 15
>>> from HowManyMedals import howManyMedals
>>> howManyMedals(data, 'Kjetil Andr Aamodt')
{1992: {'G': 1, 'S': 0, 'B': 1}, 1994: {'G': 0, 'S': 2, 'B': 1}, 1998: {'G': 0, 'S': 0, 'B':
↪  0}, 2002: {'G': 2, 'S': 0, 'B': 0}, 2006: {'G': 1, 'S': 0, 'B': 0}}
```

Exercise 04 - SpatioTemporalData

Turn-in directory :	ex04
Files to turn in :	FileLoader.py, SpatioTemporalData.py
Allowed libraries :	Pandas
Remarks :	n/a

This exercise uses the dataset `athlete_events.csv`

Write a class called `SpatioTemporalData` which takes a dataset (pandas DataFrame) as argument in its constructor and implements the following methods:

- `when(location)` : takes a location as an argument and returns a list containing the years where games were held in the given location.

- `where(date)` : takes a date as an argument and returns the location where the Olympics took place in the given year.

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
Loading dataset of dimensions 271116 x 15
>>> from SpatioTemporalData import SpatioTemporalData
>>> sp = SpatioTemporalData(data)
>>> sp.where(1896)
['Athina']
>>> sp.where(2016)
['Rio de Janeiro']
>>> sp.when('Athina')
[2004, 1906, 1896]
>>> sp.when('Paris')
[1900, 1924]
```

Exercise 05 - HowManyMedalsByCountry

Turn-in directory :	ex05
Files to turn in :	FileLoader.py, HowManyMedalsByCountry.py
Allowed libraries :	Pandas
Remarks :	n/a

This exercise uses the following dataset: `athlete_events.csv`

Write a function `howManyMedalsByCountry` which takes two arguments:

- a pandas.DataFrame which contains the dataset

- a country name

The function returns a dictionary of dictionaries giving the number and type of medal for each competition where the country team earned medals.
The keys of the main dictionary are the Olympic games' years. In each year's dictionary, the key are 'G', 'S', 'B' corresponding to the type of medals won.

Duplicated medals per team games should be handled and not counted twice.

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
Loading dataset of dimensions 271116 x 15
>>> from HowManyMedalsByCountry import howManyMedalsByCountry
>>> howManyMedalsByCountry(data, 'Martian Federation')
{2192: {'G': 17, 'S': 14, 'B': 23}, 2196: {'G': 8, 'S': 21, 'B': 19}, 2200: {'G': 26, 'S':
↪  19, 'B': 7}}
```

You probably guessed by now that we gave up providing real examples...

If you want real examples, you can easily look online. Do beware that some medals might be awarded or removed years after the games are over, for example if a previous medallist was found to have cheated and is sanctioned. The `athlete_events.csv` dataset might not always take these posterior changes into account.

Exercise 06 - MyPlotLib

Turn-in directory :	ex06
Files to turn in :	MyPlotLib.py
Allowed libraries :	Pandas, Matplotlib, Seaborn, Scipy
Remarks :	The less work you do, the better! You don't necessarily need all those libraries to complete the exercise.

This exercise uses the following dataset: `athlete_events.csv`

Write a class called `MyPlotLib`. This class implements different plotting methods, each of which take two arguments:

- a pandas.DataFrame which contains the dataset
- a list of feature names

Hint: What is a feature? https://towardsdatascience.com/feature-engineering-for-machine-learning-3a5e293a5114

- `histogram(data, features)` : plots one histogram for each numerical feature in the list
- `density(data, features)` : plots the density curve of each numerical feature in the list
- `pair_plot(data, features)` : plots a matrix of subplots (also called scatter plot matrix). On each subplot shows a scatter plot of one numerical variable against another one. The main diagonal of this matrix shows simple histograms.
- `box_plot(data, features)` : displays a box plot for each numerical variable in the dataset.

Examples

Figure 1: histogram

Remember: The less work you do, the better!

Figure 2: density

Figure 3: pair_plot

Figure 4: box_plot

Exercise 07 - Komparator

Turn-in directory :	ex07
Files to turn in :	Komparator.py, MyPlotLib.py (optional)
Allowed libraries :	Pandas, Matplotlib, Seaborn, Scipy
Remarks :	The less work you do, the better! You don't necessarily need all those libraries to complete the exercise.

This exercise uses the following dataset: `athlete_events.csv`

Write a class called `Komparator` whose constructor takes as an argument a pandas.DataFrame which contains the dataset. The class must implement the following methods, which take as input two variable names:

- `compare_box_plots(categorical_var, numerical_var)` : displays a series of box plots to compare how the distribution of the numerical variable changes if we only consider the subpopulation which belongs to each category. There should be as many box plots as categories. For example, with Sex and Height, we would compare the height distributions of men vs. women with two box plots.

- `density(categorical_var, numerical_var)` : displays the density of the numerical variable. Each subpopulation should be represented by a separate curve on the graph.

- `compare_histograms(categorical_var, numerical_var)` : plots the numerical variable in a separate histogram for each category. As a bonus, you can use overlapping histograms with a color code.

`BONUS`: Your functions can also accept a list of numerical variables (instead of just one), and output a comparison plot for each variable in the list.