

Bootcamp Python



Module04 Pandas

Module04 - Pandas

Today you will learn how to use a Python library that will allow you to manipulate dataframes.

Notions of the module

Pandas! And Bamboos! (DataFrames, Series, slicing, mask ...)

General rules

- Use the Pandas Library.
- The version of Python recommended to use is 3.7, you can check the version of Python with the following command: `python -V`
- The norm: during this bootcamp you will follow the [PEP 8 standards](#). You can install [pycodestyle](#) which is a tool to check your Python code.
- The function eval is never allowed.
- The exercises are ordered from the easiest to the hardest.
- Your exercises are going to be evaluated by someone else, so make sure that your variable names and function names are appropriate and civil.
- Your manual is the internet.
- You can also ask questions in the [#bootcamps](#) channel in the 42 AI Slack: [42-ai.slack.com](#).
- If you find any issue or mistakes in the subject please create an issue on our [bootcamp python repository on Github](#).

Helper

For this module you will use the dataset `athlete_events.csv` provided in the `resources` folder.

You will need to install pandas library:

```
pip install pandas
```

Ensure that you have the right Python version.

```
> which python
/goinfre/miniconda/bin/python
> python -V
Python 3.7.*
> which pip
/goinfre/miniconda/bin/pip
```

Exercise 00 - FileLoader

Exercise 01 - YoungestFellah

Exercise 02 - ProportionBySport

Exercise 03 - HowManyMedals

Exercise 04 - SpatioTemporalData

Exercise 05 - HowManyMedalsByCountry

Exercise 06 - MyPlotLib

Exercise 07 - Komparator

\clearpage

Exercise 00 - FileLoader

Turn-in directory:	ex00/
Files to turn in:	FileLoader.py
Allowed libraries:	Pandas
Remarks:	Be as lazy as possible...

Objective:

Basic usage of Pandas to retrieve a dataset and store it in a `PandaDataFrame`.

Instructions:

Write a class named `FileLoader` which implements the following methods:

- `load(path)` : takes as an argument the file path of the dataset to load, displays a message specifying the dimensions of the dataset (e.g. 340 x 500) and returns the dataset loaded as a `pandas.DataFrame`,
- `display(df, n)` : takes a `pandas.DataFrame` and an integer as arguments, displays the first n rows of the dataset if n is positive, or the last n rows if n is negative.

`FileLoader` object should not raise any exceptions (i.e. you should handle wrong path, file does not exist, parameters different than the expected type ...).

Examples:

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load("../data/adult_data.csv")
Loading dataset of dimensions 32561 x 15
>>> loader.display(data, 12)
```

	age	workclass	fnlwgt	...	hours-per-week	native-country	salary
0	39	State-gov	77516	...	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	...	13	United-States	<=50K
2	38	Private	215646	...	40	United-States	<=50K
3	53	Private	234721	...	40	United-States	<=50K
4	28	Private	338409	...	40	Cuba	<=50K
5	37	Private	284582	...	40	United-States	<=50K
6	49	Private	160187	...	16	Jamaica	<=50K
7	52	Self-emp-not-inc	209642	...	45	United-States	>50K

```

8      31      Private    45781 ...      50    United-States    >50K
9      42      Private   159449 ...      40    United-States    >50K
10     37      Private   280464 ...      80    United-States    >50K
11     30    State-gov   141297 ...      40           India    >50K

[12 rows x 15 columns]

```

NB: Your terminal may display more columns if the window is wider.

\clearpage

Exercise 01 - YoungestFellah

Turn-in directory:	ex01/
Files to turn in:	FileLoader.py, YoungestFellah.py
Allowed libraries:	Pandas
Remarks:	n/a

Objective:

Discover of the concept of slicing and indexing in `Pandas.DataFrame`.

Instructions:

This exercise uses the following dataset: `athlete_events.csv`

Write a function `youngestFellah` that takes two arguments as describe by the function's signature:

```

def youngestFellah(df, year):
    """
    Get the name of the youngest woman and man for the given year.
    Args:
    ----
        df: pandas.DataFrame object containing the dataset.
        year: integer corresponding to a year.
    Return:
    ----
        dct: dictionary with 2 keys for female and male athlete.
    Raises:
    ----
        This function should not raise any Exceptions.
    """
    ... your code ...

```

The function returns a dictionary containing the age of the youngest woman and man who took part in the Olympics on that year. The name of the dictionary's keys is up to you, but it must be self-explanatory.

Examples:

```

>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../ressources/athlete_events.csv')
# Output
Loading dataset of dimensions 271116 x 15

```

```
>>> from YoungestFellah import youngestFellah
>>> youngestFellah(data, 2004)
# Output
{'f': 13.0, 'm': 14.0}

>>> youngestFellah(data, 1991)
# Output
{'f': 'nan', 'm': 'nan'}
```

\clearpage

Exercise 02 - ProportionBySport

Turn-in directory:	ex02/
Files to turn in:	FileLoader.py, ProportionBySport.py
Allowed libraries:	Pandas
Remarks:	n/a

Objective:

Practise of the concept of slicing and indexing in `Pandas.DataFrame`.

Instructions:

This exercise uses the dataset `athlete_events.csv`. Write a function `proportionBySport` that takes four arguments as describe by the function's signature:

```
def proportionBySport(df, year, sport, gender):
    """
    Calculates the percentage of participants of specific gender for a specific sport among
    all the participants of the same gender for the given year.
    Args:
    -----
        df: pandas.DataFrame object containing the dataset.
        year: integer corresponding to a year.
        sport: string corresponding to a sport.
        gender: string corresponding to a gender.
    Return:
    -----
        dct: dictionary with 2 keys for female and male athlete.
    Raises:
    -----
        This function should not raise any Exceptions.
    """
    ... your code ...
```

The function returns a float corresponding to the proportion (percentage) of participants who played the given sport among the participants of the given gender.

The function answers questions like the following : “What was the percentage of female basketball players among all the female participants of the 2016 Olympics?”

Hint: Here and further, if needed, drop duplicated sports people to count only unique ones. Beware to call the dropping function at the right moment and with the right parameters, in order to not omit any individuals.

Examples:

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
# Output
Loading dataset of dimensions 271116 x 15

>>> from ProportionBySport import proportionBySport
>>> proportionBySport(data, 2004, 'Tennis', 'F')
# Output
0.01935634328358209
```

We assume that we are always using appropriate arguments as input, and thus do not need to handle input errors.

\clearpage

Exercise 3 - HowManyMedals

Turn-in directory:	ex03/
Files to turn in:	FileLoader.py, HowManyMedals.py
Allowed libraries:	Pandas
Remarks:	n/a

Objective:

Deepen the manipulation of the concept of slicing and indexing in `Pandas.DataFrame`.

Instructions:

This exercise uses the dataset: `athlete_events.csv`

Write a function `howManyMedals` that takes two arguments, as describe by the function signature:

```
def proportionBySport(df, name):
    """
    Calculates the percentage of participants of specific gender for a specific sport among
    → all the participants of the same gender for the given year.
    Args:
    -----
        df: pandas.DataFrame object containing the dataset.
        name: string corresponding to the name of a participant.
    Return:
    -----
        dct: nested dictionary {'year_1': {'G':nb_1G, 'S':nb_1S, 'B':n_1B}, 'year_2':
    → {'G':nb_2G, 'S':nb_2S, 'B':n_2B}}.
    Raises:
    -----
        This function should not raise any Exceptions.
    """
    ... your code ...
```

The function returns a dictionary of dictionaries giving the number and type of medals for each year during which the participant won medals. The keys of the main dictionary are the Olympic games years. In each year's dictionary, the keys are 'G', 'S', 'B' corresponding to the type of medals won (gold, silver, bronze). The innermost values correspond to the number of medals of a given type won for a given year.

Examples:

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../ressources/athlete_events.csv')
# Output
Loading dataset of dimensions 271116 x 15

>>> from HowManyMedals import howManyMedals
>>> howManyMedals(data, 'Kjetil Andr Aamodt')
# Output
{1992: {'G': 1, 'S': 0, 'B': 1},
 1994: {'G': 0, 'S': 2, 'B': 1},
 1998: {'G': 0, 'S': 0, 'B': 0},
 2002: {'G': 2, 'S': 0, 'B': 0},
 2006: {'G': 1, 'S': 0, 'B': 0}}
```

\clearpage

Exercise 04 - SpatioTemporalData

Turn-in directory:	ex04/
Files to turn in:	FileLoader.py, SpatioTemporalData.py
Allowed libraries:	Pandas
Remarks:	Take a look to the methods of DataFrame

Objective:

The goal of this exercise is to deepen your manipulation of `pandas.DataFrame` and its methods, especially the methods for slicing and indexing.

Instructions:

You will use the dataset `athlete_events.csv`

You have to write the class `SpatioTemporalData` which takes a dataset (`pandas.DataFrame`) as argument in its constructor and have the two following methods:

- `when(self, location)` : takes a location as an argument and returns a list containing the years where games were held in the given location.
- `where(self, date)` : takes a date as an argument and returns the location where the Olympics took place in the given year.

No exceptions should be risen by `where` and `when`. `##` Examples:

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../ressources/athlete_events.csv')
# Output
Loading dataset of dimensions 271116 x 15

>>> from SpatioTemporalData import SpatioTemporalData
>>> sp = SpatioTemporalData(data)
>>> sp.where(1896)
# Output
['Athina']
```

```
>>> sp.where(2016)
# Output
['Rio de Janeiro']

>>> sp.when('Athina')
# Output
[2004, 1906, 1896]

>>> sp.when('Paris')
# Output
[1900, 1924]
```

\clearpage

Exercise 05 - HowManyMedalsByCountry

Turn-in directory:	ex05/
Files to turn in:	FileLoader.py, HowManyMedalsByCountry.py
Allowed libraries:	Pandas
Remarks:	n/a

Objective:

The goal of this exercise is to deepen your manipulation of `pandas.DataFrame` and its methods, especially the methods for slicing and indexing.

Instructions:

You will use the following dataset: `athlete_events.csv`

You have to write a function `howManyMedalsByCountry` that takes two arguments:

- a `pandas.DataFrame` which contains the dataset
- a country name.

The function returns a nested dictionary which gives the number and type of medal for each competition where the country delegation earned medals.

The keys of the main dictionary are the Olympic games' years. In each year's dictionary, the keys are 'G', 'S', 'B' corresponding to the type of medals won.

Duplicated medals per team games should be handled and not counted twice. Plus, function should not raise any exceptions.

Examples:

```
>>> from FileLoader import FileLoader
>>> loader = FileLoader()
>>> data = loader.load('../data/athlete_events.csv')
# Output
Loading dataset of dimensions 271116 x 15

>>> from HowManyMedalsByCountry import howManyMedalsByCountry
```



```
>>> howManyMedalsByCountry(data, 'Martian Federation')
# Output
{2192: {'G': 17, 'S': 14, 'B': 23}, 2196: {'G': 8, 'S': 21, 'B': 19}, 2200: {'G': 26, 'S': 19, 'B': 7}}
```

You probably guessed but it is not a real examples.

If you want real examples, you can easily look online. Beware that some medals might be awarded or removed years after the games are over, for example if a previous medallist was found to have cheated and is sanctioned. The `athlete_events.csv` dataset might not always take these posterior changes into account.

\clearpage

Exercise 06 - MyPlotLib

Turn-in directory:	ex06/
Files to turn in:	MyPlotLib.py
Allowed libraries:	Pandas, Matplotlib, Seaborn, Scipy
Remarks:	The less work you do, the better! You don't necessarily need all those libraries to complete the exercise.

Objective:

The goal of the exercise is to introduce plotting methods among the different graphical libraries (or graphical modules within libraries) Pandas, Matplotlib, Seaborn, Scipy ...

Instructions:

You will use the dataset: `athlete_events.csv`.


You have to write a class called `MyPlotLib`. This class implements different plotting methods, each taking two arguments:

- a `pandas.DataFrame` which contains the dataset
- a list of feature (also called independant variable) names.


Here the list of methods you will implement:

- `histogram(data, features)`: plots one histogram for each numerical feature in the list.
- `density(data, features)`: plots the density curve of each numerical feature in the list.
- `pair_plot(data, features)`: plots a matrix of subplots (also called scatter plot matrix). On each subplot shows a scatter plot of one numerical variable against another one. The main diagonal of this matrix shows simple histograms.
- `box_plot(data, features)`: displays a box plot for each numerical variable in the dataset.

Examples:



histogram



pair_plot

\clearpage

Exercise 07 - Komparator

Turn-in directory:	ex07/
Files to turn in:	Komparator.py, MyPlotLib.py (optional)
Allowed libraries:	Pandas, Matplotlib, Seaborn, Scipy
Remarks:	The less work you do, the better! You don't necessarily need all those libraries to complete the exercise.

Objective:

The goal of the exercise is to introduce plotting methods among the different graphical libraries **Pandas**, **Matplotlib**, **Seaborn**, **Scipy** ...

Instructions:

You will use the dataset: `athlete_events.csv`

You have to write a class called `Komparator` whose constructor takes `pandas.DataFrame` as argument which contains the dataset. The class must implement the following methods, which take as input two variable names:

- `compare_box_plots(categorical_var, numerical_var)` : displays box plot with several boxes which allow to compare the distribution of the numerical variable according to the categorical values of `categorical_var`. There should be as many box as categories. For example, with Sex and Height, we would compare the height distributions of men vs. women via two boxes.
- `density(categorical_var, numerical_var)` : displays the density of the numerical variable. Each subpopulation should be represented by a separate curve on the graph.
- `compare_histograms(categorical_var, numerical_var)` : plots the numerical variable in a separate histogram for each category. As a bonus, you can use overlapping histograms with a color code.

Bonus: Your functions can also accept a list of numerical variables (instead of just one), and output a comparison plot for each variable in the list.

\clearpage